

PSTAT131_HW1

Justin Lee

10/14/2020

(1)

Reading the data and attaching packages

```
library(tidyverse)
library(dplyr)
library(ggplot2)
library(resample)
library(miscTools)
library(FIACH)

algae <- read_table2("algaeBloom.txt", col_names=c('season', 'size', 'speed', 'mxPH', 'mn02', 'Cl', 'N03', 'NH4',
'oP04', 'P04', 'Chla', 'a1', 'a2', 'a3', 'a4', 'a5', 'a6', 'a7'),
na="XXXXXXX")

glimpse(algae)
```

```
## Rows: 200
## Columns: 18
## $ season <chr> "winter", "spring", "autumn", "spring", "autumn", "winter", ...
## $ size <chr> "small", "small", "small", "small", "small", "small", "small..."
## $ speed <chr> "medium", "medium", "medium", "medium", "medium", "high", "h..."
## $ mxPH <dbl> 8.00, 8.35, 8.10, 8.07, 8.06, 8.25, 8.15, 8.05, 8.70, 7.93, ...
## $ mn02 <dbl> 9.8, 8.0, 11.4, 4.8, 9.0, 13.1, 10.3, 10.6, 3.4, 9.9, 10.2, ...
## $ Cl <dbl> 60.800, 57.750, 40.020, 77.364, 55.350, 65.750, 73.250, 59.0...
## $ N03 <dbl> 6.238, 1.288, 5.330, 2.302, 10.416, 9.248, 1.535, 4.990, 0.8...
## $ NH4 <dbl> 578.000, 370.000, 346.667, 98.182, 233.700, 430.000, 110.000...
## $ oP04 <dbl> 105.000, 428.750, 125.667, 61.182, 58.222, 18.250, 61.250, 4...
## $ P04 <dbl> 170.000, 558.750, 187.057, 138.700, 97.580, 56.667, 111.750, ...
## $ Chla <dbl> 50.000, 1.300, 15.600, 1.400, 10.500, 28.400, 3.200, 6.900, ...
## $ a1 <dbl> 0.0, 1.4, 3.3, 3.1, 9.2, 15.1, 2.4, 18.2, 25.4, 17.0, 16.6, ...
## $ a2 <dbl> 0.0, 7.6, 53.6, 41.0, 2.9, 14.6, 1.2, 1.6, 5.4, 0.0, 0.0, 0....
## $ a3 <dbl> 0.0, 4.8, 1.9, 18.9, 7.5, 1.4, 3.2, 0.0, 2.5, 0.0, 0.0, 0.0, ...
## $ a4 <dbl> 0.0, 1.9, 0.0, 0.0, 0.0, 0.0, 3.9, 0.0, 0.0, 2.9, 0.0, 0.0, ...
## $ a5 <dbl> 34.2, 6.7, 0.0, 1.4, 7.5, 22.5, 5.8, 5.5, 0.0, 0.0, 1.2, 0.0...
## $ a6 <dbl> 8.3, 0.0, 0.0, 0.0, 4.1, 12.6, 6.8, 8.7, 0.0, 0.0, 0.0, 0.0, ...
## $ a7 <dbl> 0.0, 2.1, 9.7, 1.4, 1.0, 2.9, 0.0, 0.0, 0.0, 1.7, 6.0, 1.5, ...
```

(a)

```
algae %>% group_by(season) %>% summarise(n = n())
```

```
## # A tibble: 4 x 2
##   season      n
##   <chr>   <int>
## 1 autumn    40
## 2 spring    53
## 3 summer    45
## 4 winter    62
```

From the data we see above, we can see the total count of the observations are:

- Autumn = 40
- Spring = 53
- Summer = 45
- Winter = 62

(b)

```
missing = is.na(algae)
length(missing[missing == TRUE])
```

```
## [1] 33
```

```
chemicals = algae[, 4:11]
colMeans(chemicals, na.rm = TRUE)
```

```
##      mxPH      mnO2      Cl      NO3      NH4      oP04      P04
##  8.011734  9.117778 43.636279  3.282389 501.295828  73.590596 137.882101
##      Chla
## 13.971197
```

```
#Using package 'Resample'
colVars(chemicals, na.rm = TRUE)
```

```
##      mxPH      mnO2      Cl      NO3      NH4      oP04
## 3.579693e-01 5.718089e+00 2.193172e+03 1.426176e+01 3.851585e+06 8.305850e+03
##      P04      Chla
## 1.663938e+04 4.200827e+02
```

We can confirm that there are missing values within the dataset using the function “is.na()”. The total count of the missing values were 33. I used the length function to see how many counts of “TRUE” there were in the “missing” subset. The mean and variance is shown above as well using the *colMeans* function as well as *colvars* from a package ‘Resample’ that I have learned from another class. For both the cases of mean and variance, the missing values were ignored. The magnitude is greater than the mean for most of the chemicals which can indicate that the data points are very spread out from the average.

(c)

```
colMedians(chemicals, na.rm = TRUE)
```

```
##      mxPH      mnO2      Cl      NO3      NH4      oP04      P04      Chla
##    8.0600    9.8000   32.7300    2.6750  103.1665   40.1500  103.2855    5.4750
```

```
#chemicals subset with the missing values replaced with mean for each chemical
chemicals_2 = chemicals
chemicals_2$mnO2[is.na(chemicals_2$mnO2)] = mean(chemicals$mnO2,na.rm = TRUE)
chemicals_2$Cl[is.na(chemicals_2$Cl)] = mean(chemicals$Cl,na.rm = TRUE)
chemicals_2$NO3[is.na(chemicals_2$NO3)] = mean(chemicals$NO3,na.rm = TRUE)
chemicals_2$NH4[is.na(chemicals_2$NH4)] = mean(chemicals$NH4,na.rm = TRUE)
chemicals_2$oP04[is.na(chemicals_2$oP04)] = mean(chemicals$oP04,na.rm = TRUE)
chemicals_2$P04[is.na(chemicals_2$P04)] = mean(chemicals$P04,na.rm = TRUE)
chemicals_2$Chla[is.na(chemicals_2$Chla)] = mean(chemicals$Chla,na.rm = TRUE)

missing2 = is.na(chemicals_2)
#Confirming there are no missing values, this should output zero if I did everythin correctly
length(missing2[missing2 == TRUE])
```

```
## [1] 1
```

```
#Computing M.A.D. for each chemical using package 'FIACH'
colMad(chemicals_2)
```

```
## [1] 0.504084 1.979271 35.337771 2.153477 112.059356 45.466153 121.444955
## [8] 7.685057
```

Since the output for the Median Absolute Deviation is does not show which output for each chemical, I will state them neatly here:

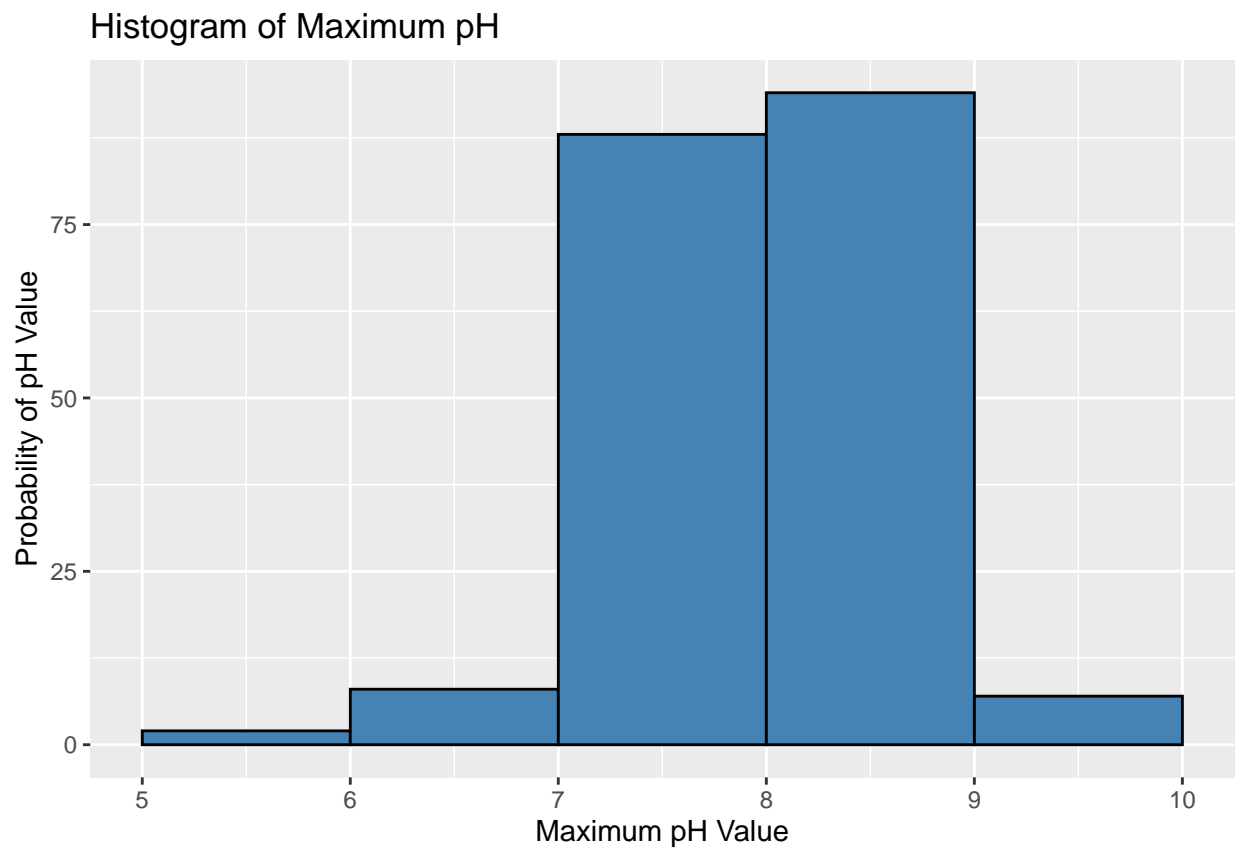
- mxPH = 0.504084
- mnO2 = 1.979271
- Cl = 35.337771
- NO3 = 2.153477
- NH4 = 112.059356
- OPO4 = 45.466153
- PO4 = 121.444955
- Chla = 7.685057

The median for each chemicals were found using the *colMedians* function from the package “miscTools”. This provided a lot more of a simple way for me to output the name of the chemical and the median. Finding the median absolute deviation has many ways including using the *mad()* function but for simplicity I used *colMad* from the package “FIACH” that I learned in the past. Since missing values present needs to be replaced to find M.A.D., I used *is.na()* function to find the missing values and replaced them with the mean value for that chemical. This adjusted value should remain closer to the original, but just more accurate. We can see that the mean and variances have magnitudes with larger differences. We can see that the variance for *mnO₂* is 5.718089 and the mean is 9.117778. This is almost two times from the variance and there are significant differences between chemicals in variance and mean. However, the medians and the median absolute deviations of the chemicals are very close to each other compared to the other. This pattern is true for all of the chemicals beside *mnO₂*.

(2)

(a)

```
algae %>% ggplot(aes(x=mxPH, stat = "density")) +  
  geom_histogram(breaks = seq(5, 10, by = 1), col = "black", fill = "steelblue") +  
  labs(title = "Histogram of Maximum pH", x = "Maximum pH Value", y = "Probability of pH Value")  
  
## Warning: Removed 1 rows containing non-finite values (stat_bin).
```




I used the `ggplot()` function to produce a histogram with probability on the vertical axis and the maximum pH on the horizontal axis. Using the statement `stat = "density"` gives us a histogram that contains a measure of density instead of frequency. The distribution seems to be skewed slightly to the left from the plot and the computed median of `mxPH` is larger than it's mean.

(b)

```
algae %>% ggplot(aes(x=mxPH, stat = "density")) +  
  geom_histogram(breaks = seq(5, 10, by = 1), col = "black", fill = "steelblue") +  
  geom_density(aes(y = ..density..*(100))) +  
  geom_rug(col = "red") +  
  labs(title = "Histogram of Maximum pH", x = "Maximum pH Value", y = "Probability of pH Value") +  
  ylim(c(0,100))
```

```
## Warning: Removed 1 rows containing non-finite values (stat_bin).
```

```
## Warning: Removed 1 rows containing non-finite values (stat_density).
```




Using `geom_density()` and `geom_rug`, we can add the density curve along with the rug plot. The `geom_density()` did give me an error in the beginning due to the fact that the probability was out of a hundred, but the density curve was out of one. However, I multiplied the y-values of the density curve by a hundred to fix the problem.

(c)

```
algae %>% ggplot(aes(x = size, y = a1)) +  
  geom_boxplot(col = "steelblue") +  
  labs(title = expression(paste("A conditioned Boxplot of Algal ", "a"[1])), x = "Size of River", y = e
```

A conditioned Boxplot of Algal a_1



We use `ggplot()` along with `geom_boxplot` to create a boxplot for a_1 . We use the `aes()` statement to group them by size which are: small,medium,large. We also specify in the `aes()` statement to indicate that the y-axis will be the data from a_1 and the x-axis will the size.

(d)

```
x <- algae$NO3
y <- algae$NH4
x[which(x %in% boxplot.stats(x)$out)]
```

```
## [1] 10.416  9.248  9.773  9.715 45.650
```

```
y[which(y %in% boxplot.stats(y)$out)]
```

```
## [1] 578.000 8777.600 1729.000 3515.000 6400.000 1911.000 647.570
## [8] 1386.250 2082.850 2167.370 737.500 914.000 5738.330 4073.330
## [15] 758.750 931.833 723.667 3466.660 920.000 1990.160 24064.000
## [22] 1131.660 1495.000 643.000 627.273 1168.000 1081.660
```

To find the outliers we can use the range given by this formula: $[(Q1 - 1.5IQR), (Q3 + 1.5IQR)]$. IQR is the interquartile range and Q1 and Q3 are the quartiles. I assigned the variable NO_3 to `x` and NH_4 to `y`. Using the operator, `which()`, and the function, `boxplot.stats()`. Through the function I wrote, we are able to find which data points are outliers. From the results we can see that there are 5 outliers for NO_3 and 27 outliers for NH_4 .

(e)

From question 1 we know:

NO_3 :

mean = 3.282389, variance = 14.26176, median = 2.6750, MAD = 22.153477

NH_4 :

mean = 501.295828, variance = 3851585, median = 103.1665, MAD = 112.059356

We can see that both the variance of NO_3 and NH_4 is significantly larger than the mean. When we look at the values for the median and MAD we can see that they are very similar. In conclusion, the median and MAD values that are not influenced much from the outlier which means that they are more robust measurements than the mean and variance.

(3)

(a)

```
summary(algae)
```

```
##      season      size      speed      mxPH
## Length:200    Length:200    Length:200    Min.   :5.600
## Class :character Class :character Class :character 1st Qu.:7.700
## Mode  :character Mode  :character Mode  :character Median :8.060
##                                     Mean  :8.012
##                                     3rd Qu.:8.400
##                                     Max.   :9.700
##                                     NA's    :1
##      mn02      C1      NO3      NH4
## Min.   : 1.500 Min.   : 0.222 Min.   : 0.050 Min.   : 5.00
## 1st Qu.: 7.725 1st Qu.: 10.981 1st Qu.: 1.296 1st Qu.: 38.33
## Median : 9.800 Median : 32.730 Median : 2.675 Median : 103.17
## Mean   : 9.118 Mean   : 43.636 Mean   : 3.282 Mean   : 501.30
## 3rd Qu.:10.800 3rd Qu.: 57.824 3rd Qu.: 4.446 3rd Qu.: 226.95
## Max.   :13.400 Max.   :391.500 Max.   :45.650 Max.   :24064.00
## NA's    :2    NA's    :10    NA's    :2    NA's    :2
##      oP04      P04      Chla      a1
## Min.   : 1.00 Min.   : 1.00 Min.   : 0.200 Min.   : 0.00
## 1st Qu.: 15.70 1st Qu.: 41.38 1st Qu.: 2.000 1st Qu.: 1.50
## Median : 40.15 Median :103.29 Median : 5.475 Median : 6.95
## Mean   : 73.59 Mean   :137.88 Mean   :13.971 Mean   :16.92
## 3rd Qu.: 99.33 3rd Qu.:213.75 3rd Qu.:18.308 3rd Qu.:24.80
## Max.   :564.60 Max.   :771.60 Max.   :110.456 Max.   :89.80
## NA's    :2    NA's    :2    NA's    :12
##      a2      a3      a4      a5
## Min.   : 0.000 Min.   : 0.000 Min.   : 0.000 Min.   : 0.000
## 1st Qu.: 0.000 1st Qu.: 0.000 1st Qu.: 0.000 1st Qu.: 0.000
## Median : 3.000 Median : 1.550 Median : 0.000 Median : 1.900
## Mean   : 7.458 Mean   : 4.309 Mean   : 1.992 Mean   : 5.064
## 3rd Qu.:11.375 3rd Qu.: 4.925 3rd Qu.: 2.400 3rd Qu.: 7.500
## Max.   :72.600 Max.   :42.800 Max.   :44.600 Max.   :44.400
##
##      a6      a7
## Min.   : 0.000 Min.   : 0.000
## 1st Qu.: 0.000 1st Qu.: 0.000
## Median : 0.000 Median : 1.000
## Mean   : 5.964 Mean   : 2.495
## 3rd Qu.: 6.925 3rd Qu.: 2.400
## Max.   :77.600 Max.   :31.600
##
```

```
sum(is.na(algae))
```

```
## [1] 33
```


Using the `summary()` function, we can see which variables have missing values and how many there are. We can see that every predictors have missing values starting from `mxPH` to `Chla`. From this we can see that `mxPH` has 1 missing value, `mnO2` has 2 missing values, `Cl` has 10 missing values, `NO3` has 2 missing values, `NH4` has 2 missing values, `OPO4` has 2 missing values, `PO4` has 2 missing values, and `Chla` has 12 missing values. This brings us to a total of 33 missing values.

(b)

```
algae.del = filter(algae, !is.na(mxPH), !is.na(mnO2), !is.na(Cl), !is.na(NO3), !is.na(NH4), !is.na(oPO4))
summary(algae.del)
```

```
##      season      size      speed      mxPH
## Length:184    Length:184    Length:184    Min.   :7.000
## Class :character Class :character Class :character 1st Qu.:7.777
## Mode  :character Mode  :character Mode  :character Median :8.100
##                                     Mean  :8.078
##                                     3rd Qu.:8.400
##                                     Max.   :9.500
##
##      mnO2      Cl      NO3      NH4
## Min.   : 1.500    Min.   : 0.80    Min.   : 0.050    Min.   : 5.80
## 1st Qu.: 7.675    1st Qu.: 11.85    1st Qu.: 1.364    1st Qu.: 49.38
## Median : 9.750    Median : 35.08    Median : 2.820    Median : 115.71
## Mean   : 9.019    Mean   : 44.88    Mean   : 3.384    Mean   : 537.67
## 3rd Qu.:10.700    3rd Qu.: 58.52    3rd Qu.: 4.540    3rd Qu.: 235.25
## Max.   :13.400    Max.   :391.50    Max.   :45.650    Max.   :24064.00
##
##      oPO4      PO4      Chla      a1
## Min.   : 1.25    Min.   : 2.50    Min.   : 0.200    Min.   : 0.00
## 1st Qu.: 18.56    1st Qu.: 50.34    1st Qu.: 2.075    1st Qu.: 1.40
## Median : 46.28    Median :115.60    Median : 5.522    Median : 4.85
## Mean   : 78.27    Mean   :146.58    Mean   : 13.883    Mean   :15.32
## 3rd Qu.:102.83    3rd Qu.:220.25    3rd Qu.: 18.308    3rd Qu.:19.32
## Max.   :564.60    Max.   :771.60    Max.   :110.456    Max.   :89.80
##
##      a2      a3      a4      a5
## Min.   : 0.000    Min.   : 0.000    Min.   : 0.000    Min.   : 0.000
## 1st Qu.: 0.000    1st Qu.: 0.000    1st Qu.: 0.000    1st Qu.: 0.000
## Median : 3.600    Median : 1.700    Median : 0.000    Median : 2.650
## Mean   : 7.777    Mean   : 4.613    Mean   : 1.846    Mean   : 5.493
## 3rd Qu.:11.700    3rd Qu.: 5.525    3rd Qu.: 2.425    3rd Qu.: 8.000
## Max.   :72.600    Max.   :42.800    Max.   :44.600    Max.   :44.400
##
##      a6      a7
## Min.   : 0.000    Min.   : 0.000
## 1st Qu.: 0.000    1st Qu.: 0.000
## Median : 0.000    Median : 1.000
## Mean   : 6.447    Mean   : 2.665
## 3rd Qu.: 7.975    3rd Qu.: 2.700
## Max.   :77.600    Max.   :31.600
```

```
str(algae.del)
```

```
## tibble [184 x 18] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
## $ season: chr [1:184] "winter" "spring" "autumn" "spring" ...
```

```
## $ size : chr [1:184] "small" "small" "small" "small" ...
## $ speed : chr [1:184] "medium" "medium" "medium" "medium" ...
## $ mxPH : num [1:184] 8 8.35 8.1 8.07 8.06 8.25 8.15 8.05 8.7 7.93 ...
## $ mnO2 : num [1:184] 9.8 8 11.4 4.8 9 13.1 10.3 10.6 3.4 9.9 ...
## $ Cl : num [1:184] 60.8 57.8 40 77.4 55.4 ...
## $ NO3 : num [1:184] 6.24 1.29 5.33 2.3 10.42 ...
## $ NH4 : num [1:184] 578 370 346.7 98.2 233.7 ...
## $ oPO4 : num [1:184] 105 428.8 125.7 61.2 58.2 ...
## $ PO4 : num [1:184] 170 558.8 187.1 138.7 97.6 ...
## $ Chla : num [1:184] 50 1.3 15.6 1.4 10.5 ...
## $ a1 : num [1:184] 0 1.4 3.3 3.1 9.2 15.1 2.4 18.2 25.4 17 ...
## $ a2 : num [1:184] 0 7.6 53.6 41 2.9 14.6 1.2 1.6 5.4 0 ...
## $ a3 : num [1:184] 0 4.8 1.9 18.9 7.5 1.4 3.2 0 2.5 0 ...
## $ a4 : num [1:184] 0 1.9 0 0 0 0 3.9 0 0 2.9 ...
## $ a5 : num [1:184] 34.2 6.7 0 1.4 7.5 22.5 5.8 5.5 0 0 ...
## $ a6 : num [1:184] 8.3 0 0 0 4.1 12.6 6.8 8.7 0 0 ...
## $ a7 : num [1:184] 0 2.1 9.7 1.4 1 2.9 0 0 0 1.7 ...
## - attr(*, "spec")=
## .. cols(
## .. season = col_character(),
## .. size = col_character(),
## .. speed = col_character(),
## .. mxPH = col_double(),
## .. mnO2 = col_double(),
## .. Cl = col_double(),
## .. NO3 = col_double(),
## .. NH4 = col_double(),
## .. oPO4 = col_double(),
## .. PO4 = col_double(),
## .. Chla = col_double(),
## .. a1 = col_double(),
## .. a2 = col_double(),
## .. a3 = col_double(),
## .. a4 = col_double(),
## .. a5 = col_double(),
## .. a6 = col_double(),
## .. a7 = col_double()
## .. )
```

There are a total of *184 observations* in `algae.del`.

(c)

```
algae.med = algae %>%
mutate_at(vars(mxPH,mnO2,Cl,NO3,NH4,oPO4,PO4,Chla),
  funs(ifelse(is.na(.)==TRUE,median(algae$.,na.rm = TRUE),.)))
```

```
## Warning: funs() is soft deprecated as of dplyr 0.8.0
## Please use a list of either functions or lambdas:
##
## # Simple named list:
```

```
## list(mean = mean, median = median)
##
## # Auto named with `tibble::lst()`
## tibble::lst(mean, median)
##
## # Using lambdas
## list(~ mean(., trim = .2), ~ median(., na.rm = TRUE))
## This warning is displayed once per session.
```

```
str(algae.med)
```

```
## tibble [200 x 18] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
## $ season: chr [1:200] "winter" "spring" "autumn" "spring" ...
## $ size : chr [1:200] "small" "small" "small" "small" ...
## $ speed : chr [1:200] "medium" "medium" "medium" "medium" ...
## $ mxPH : num [1:200] 8 8.35 8.1 8.07 8.06 8.25 8.15 8.05 8.7 7.93 ...
## $ mnO2 : num [1:200] 9.8 8 11.4 4.8 9 13.1 10.3 10.6 3.4 9.9 ...
## $ Cl : num [1:200] 60.8 57.8 40 77.4 55.4 ...
## $ NO3 : num [1:200] 6.24 1.29 5.33 2.3 10.42 ...
## $ NH4 : num [1:200] 578 370 346.7 98.2 233.7 ...
## $ oP04 : num [1:200] 105 428.8 125.7 61.2 58.2 ...
## $ P04 : num [1:200] 170 558.8 187.1 138.7 97.6 ...
## $ Chla : num [1:200] 50 1.3 15.6 1.4 10.5 ...
## $ a1 : num [1:200] 0 1.4 3.3 3.1 9.2 15.1 2.4 18.2 25.4 17 ...
## $ a2 : num [1:200] 0 7.6 53.6 41 2.9 14.6 1.2 1.6 5.4 0 ...
## $ a3 : num [1:200] 0 4.8 1.9 18.9 7.5 1.4 3.2 0 2.5 0 ...
## $ a4 : num [1:200] 0 1.9 0 0 0 0 3.9 0 0 2.9 ...
## $ a5 : num [1:200] 34.2 6.7 0 1.4 7.5 22.5 5.8 5.5 0 0 ...
## $ a6 : num [1:200] 8.3 0 0 0 4.1 12.6 6.8 8.7 0 0 ...
## $ a7 : num [1:200] 0 2.1 9.7 1.4 1 2.9 0 0 0 1.7 ...
```

```
algae.med[48,]
```

```
## # A tibble: 1 x 18
## season size speed mxPH mnO2 Cl NO3 NH4 oP04 P04 Chla a1 a2
## <chr> <chr> <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 winter small low 8.06 12.6 9 0.23 10 5 6 1.1 35.5 0
## # ... with 5 more variables: a3 <dbl>, a4 <dbl>, a5 <dbl>, a6 <dbl>, a7 <dbl>
```

```
algae.med[62,]
```

```
## # A tibble: 1 x 18
## season size speed mxPH mnO2 Cl NO3 NH4 oP04 P04 Chla a1 a2
## <chr> <chr> <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 summer small medi~ 6.4 9.8 32.7 2.68 103. 40.2 14 5.48 19.4 0
## # ... with 5 more variables: a3 <dbl>, a4 <dbl>, a5 <dbl>, a6 <dbl>, a7 <dbl>
```

```
algae.med[199,]
```

```
## # A tibble: 1 x 18
## season size speed mxPH mnO2 Cl NO3 NH4 oP04 P04 Chla a1 a2
```

```
##    <chr>  <chr> <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 winter large medi~      8   7.6  32.7  2.68 103.  40.2 103.  5.48    0 12.5
## # ... with 5 more variables: a3 <dbl>, a4 <dbl>, a5 <dbl>, a6 <dbl>, a7 <dbl>
```

There are a total of *200 observations* in `algae.med`

(d)

```
df = data.frame(algae.del[, 4:11])

cor(df, use = "pairwise.complete.obs" )
```

```
##           mxPH      mnO2      Cl      NO3      NH4      oP04
## mxPH  1.00000000 -0.10269374  0.14709539 -0.1721302 -0.15429757  0.09022909
## mnO2 -0.10269374  1.00000000 -0.26324536  0.1179077 -0.07826816 -0.39375269
## Cl   0.14709539 -0.26324536  1.00000000  0.2109583  0.06598336  0.37925596
## NO3  -0.17213024  0.11790769  0.21095831  1.0000000  0.72467766  0.13301452
## NH4  -0.15429757 -0.07826816  0.06598336  0.7246777  1.00000000  0.21931121
## oP04  0.09022909 -0.39375269  0.37925596  0.1330145  0.21931121  1.00000000
## P04   0.10132957 -0.46396073  0.44519118  0.1570297  0.19939575  0.91196460
## Chla  0.43182377 -0.13121671  0.14295776  0.1454929  0.09120406  0.10691478
##           P04      Chla
## mxPH  0.1013296  0.43182377
## mnO2 -0.4639607 -0.13121671
## Cl   0.4451912  0.14295776
## NO3  0.1570297  0.14549290
## NH4  0.1993958  0.09120406
## oP04  0.9119646  0.10691478
## P04   1.0000000  0.24849223
## Chla  0.2484922  1.00000000
```

```
model = lm(P04~oP04, data = algae)
summary(model)
```

```
##
## Call:
## lm(formula = P04 ~ oP04, data = algae)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -110.12  -36.34  -12.68   23.26  216.98
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   42.897      4.808   8.922 3.34e-16 ***
## oP04           1.293       0.041  31.535 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 52.37 on 195 degrees of freedom
## (3 observations deleted due to missingness)
```

```
## Multiple R-squared:  0.8361, Adjusted R-squared:  0.8352
## F-statistic: 994.5 on 1 and 195 DF,  p-value: < 2.2e-16
```

```
predict(model,algae[28,"oP04"])
```

```
##          1
## 48.06929
```

```
#filling in the missing value in 'algae'
algae[28,"P04"] = predict(model,algae[28,"oP04"])
```

The value that we obtain for the missing value for PO4 based on oPO4 in the 28th observation is 48.06929.

(e)

We know from lecture that survivorship bias favors the values that appear but it ignores the values that did not appear. The bullet holes on planes example studies the concentration of the holes and the lack of bullet holes on the planes that did survive. We are not considering the fact that missing values may be indicative of an outlying phenomenon by using imputed values.

(4)

(a)

```
set.seed(123)
cv = sample(cut(1:200,breaks = 5, label = FALSE))
cv

##      [1] 4 5 1 5 5 2 3 2 5 5 4 3 3 5 5 3 4 3 2 1 1 5 5 5 2 3 5 3 3 2 4 1 3 5 5 2 1
##     [38] 4 2 4 5 5 1 2 5 5 2 4 3 5 1 1 2 5 2 1 3 3 1 3 5 1 5 5 4 1 1 4 4 2 1 3 5 3
##     [75] 1 3 1 1 3 1 1 1 4 4 2 5 5 3 3 2 2 5 2 3 1 1 2 5 5 3 5 4 5 3 2 1 2 4 2 4 3
##    [112] 1 4 5 4 5 5 4 3 1 2 1 2 4 4 3 4 2 1 4 5 1 2 2 3 1 4 5 4 4 4 2 3 5 1 2 1 4
##    [149] 1 1 4 2 4 2 2 4 3 3 1 4 4 2 2 3 3 2 3 1 2 4 1 4 3 1 5 4 3 2 5 2 2 1 2 4 4
##    [186] 4 3 5 5 4 3 1 2 3 3 3 4 3 2 1
```

(b)

```
do.chunk <- function(chunkid, chunkdef, dat){ # function argument

  train = (chunkdef != chunkid)

  Xtr = dat[train,1:11] # get training set
  Ytr = dat[train,12] # get true response values in training set

  Xvl = dat[!train,1:11] # get validation set
  Yvl = dat[!train,12] # get true response values in validation set

  lm.a1 <- lm(a1~., data = dat[train,1:12])
  predYtr = predict(lm.a1) # predict training values
  predYvl = predict(lm.a1,Xvl) # predict validation values
  data.frame(fold = chunkid,
             train.error = mean((predYtr - Ytr$a1)^2), # compute and store training error
             val.error = mean((predYvl - Yvl$a1)^2)) # compute and store test error
}

#My code
lapply(1:5, FUN = do.chunk, chunkdef = cv, dat = algae.med)

## [[1]]
##      fold train.error val.error
## 1      1      280.7503  322.0817
##
## [[2]]
##      fold train.error val.error
## 1      2      305.6518  229.1328
##
## [[3]]
##      fold train.error val.error
## 1      3      272.3432  360.8609
##
```

```
## [[4]]
##   fold train.error val.error
## 1    4    281.0626   422.842
##
## [[5]]
##   fold train.error val.error
## 1    5    270.3345   386.2931
```

(5)

```
algae.Test <- read_table2('algaeTest.txt',
                        col_names=c('season','size','speed','mxPH','mnO2','Cl','N03',
                                     'NH4','oPO4','PO4','Chla','a1'),
                        na=c('XXXXXXXX'))
```

```
## Parsed with column specification:
## cols(
##   season = col_character(),
##   size = col_character(),
##   speed = col_character(),
##   mxPH = col_double(),
##   mnO2 = col_double(),
##   Cl = col_double(),
##   N03 = col_double(),
##   NH4 = col_double(),
##   oPO4 = col_double(),
##   PO4 = col_double(),
##   Chla = col_double(),
##   a1 = col_double()
## )
```

```
firstdata = algae.med[12]
newdata = algae.Test[12]
```

```
fit = lm(a1 ~ ., data = algae.med[1:12])
```

```
firstpredict = predict(fit, algae.med[1:11])
newpredict = predict(fit, algae.Test[1:11])
```

```
data.frame(train.error = mean((firstpredict - firstdata$a1)^2), val.error = mean((newpredict - newdata$a1)^2))
```

```
##   train.error val.error
## 1    286.2661   250.1794
```

Yes, this is what is roughly expected based of the CV estimated test error from number 4. The *train.error* is 286.2661 which is very close to train.error predicted in number 4. The *val.error* is 250.1794 which is not close to the predicted val.error besides the 2nd fold.

(6)

```
library(ISLR)
```

```
## Warning: package 'ISLR' was built under R version 3.6.3
```


```
head(Wage)
```

```
##      year age      maritl      race      education      region
## 231655 2006  18 1. Never Married 1. White  1. < HS Grad 2. Middle Atlantic
## 86582  2004  24 1. Never Married 1. White  4. College Grad 2. Middle Atlantic
## 161300 2003  45      2. Married 1. White  3. Some College 2. Middle Atlantic
## 155159 2003  43      2. Married 3. Asian  4. College Grad 2. Middle Atlantic
## 11443  2005  50      4. Divorced 1. White  2. HS Grad 2. Middle Atlantic
## 376662 2008  54      2. Married 1. White  4. College Grad 2. Middle Atlantic
##
##      jobclass      health health_ins logwage      wage
## 231655 1. Industrial 1. <=Good 2. No 4.318063 75.04315
## 86582  2. Information 2. >=Very Good 2. No 4.255273 70.47602
## 161300 1. Industrial 1. <=Good 1. Yes 4.875061 130.98218
## 155159 2. Information 2. >=Very Good 1. Yes 5.041393 154.68529
## 11443  2. Information 1. <=Good 1. Yes 4.318063 75.04315
## 376662 2. Information 2. >=Very Good 1. Yes 4.845098 127.11574
```

(a)

```
Wage %>% ggplot(aes(x = age, y = wage)) +
  geom_point() +
  geom_smooth(color = "red") +
  labs(title = "Wages And Age", x = "Age", y = "Wage")
```

```
## `geom_smooth()` using method = 'gam' and formula 'y ~ s(x, bs = "cs")'
```

To plot this graph I used `ggplot()` along with `geom_point()` and `geom_smooth()`. I used the red color for the fit so that it'll stick out more. There is a pattern that we can observe and we see that the wages steadily increases until the age hits around 40 then remains constant until about age 60. Then we can see that the wage goes down slowly until the end. This matches exactly what I expected because as time goes on from you twenties, we tend to work on our skills which will give us promotions or better jobs. Then as our age passes 60, we have to start thinking about retirement or even retire.

(b)

(i):

```
attach(Wage)
x = lm(wage ~ 1 + age + I(age^2) + I(age^3) + I(age^4) + I(age^5) + I(age^6) + I(age^7) + I(age^8) + I(age^9) + I(age^10))
summary(x)

##
## Call:
## lm(formula = wage ~ 1 + age + I(age^2) + I(age^3) + I(age^4) +
##      I(age^5) + I(age^6) + I(age^7) + I(age^8) + I(age^9) + I(age^10))
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -100.38  -24.45   -4.97   15.49  199.61
##
```

```
## Coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept)  1.773e+04  2.636e+04   0.672   0.501
## age         -4.259e+03  6.821e+03  -0.624   0.532
## I(age^2)      4.412e+02  7.726e+02   0.571   0.568
## I(age^3)     -2.585e+01  5.048e+01  -0.512   0.609
## I(age^4)      9.471e-01  2.109e+00   0.449   0.653
## I(age^5)     -2.257e-02  5.894e-02  -0.383   0.702
## I(age^6)      3.513e-04  1.117e-03   0.315   0.753
## I(age^7)     -3.476e-06  1.418e-05  -0.245   0.806
## I(age^8)      2.031e-08  1.156e-07   0.176   0.860
## I(age^9)     -5.868e-11  5.468e-10  -0.107   0.915
## I(age^10)     4.646e-14  1.141e-12   0.041   0.968
##
## Residual standard error: 39.89 on 2989 degrees of freedom
## Multiple R-squared:  0.08912,    Adjusted R-squared:  0.08607
## F-statistic: 29.24 on 10 and 2989 DF,  p-value: < 2.2e-16
```

(ii):

```
set.seed(123)
library(plyr)
```

```
## Warning: package 'plyr' was built under R version 3.6.3
```

```
## -----
## You have loaded plyr after dplyr - this is likely to cause problems.
## If you need functions from both plyr and dplyr, please load plyr first, then dplyr:
## library(plyr); library(dplyr)
```

```
## -----
```

```
##
## Attaching package: 'plyr'
```

```
## The following objects are masked from 'package:dplyr':
##
##   arrange, count, desc, failwith, id, mutate, rename, summarise,
##   summarize
```

```
## The following object is masked from 'package:purrr':
##
##   compact
```

```
do.chunk_2 <- function(chunkid, chunkdef, dat, a){ # function argument

  train = (chunkdef != chunkid)

  Xtr = dat[train,1:10] # get training set
```

```

Ytr = dat[train,11] # get true response values in training set

Xvl = dat[!train,1:10] # get validation set
Yvl = dat[!train,11] # get true response values in validation set

if (a == 0){
  lm.x = lm(wage~1, data = dat[train, 1:11])
}
else {
  lm.x = lm(wage~poly(x = age, degree = a, raw = FALSE),          data = dat[train, 1:11])
}

predYtr = predict(lm.x) # predict training values
predYvl = predict(lm.x,Xvl) # predict validation values
data.frame(fold = chunkid,
            train.error = mean((predYtr - Ytr)^2), # compute and store training error
            val.error = mean((predYvl - Yvl)^2)) # compute and store test error
}

cv = sample(cut(1:nrow(Wage),breaks = 5, label = FALSE))

df <- data.frame()
for (i in 0:10){
  ldply_out <- ldply(1:5, .fun = do.chunk_2,chunkdef = cv, dat = Wage, a = i)
  df <- rbind(df, ldply_out)
}

df

```

```

##      fold train.error val.error
## 1      1    1721.594  1817.527
## 2      2    1658.553  2069.315
## 3      3    1743.628  1729.083
## 4      4    1774.148  1606.896
## 5      5    1805.085  1484.879
## 6      1    1655.070  1750.247
## 7      2    1582.324  2042.891
## 8      3    1683.087  1639.366
## 9      4    1710.395  1529.071
## 10     5    1738.332  1419.121
## 11     1    1577.534  1680.028
## 12     2    1511.439  1945.377
## 13     3    1604.291  1572.308
## 14     4    1632.043  1461.176
## 15     5    1662.475  1341.647
## 16     1    1572.410  1673.979

```

```
## 17      2      1507.355  1935.783
## 18      3      1599.535  1565.365
## 19      4      1626.202  1458.256
## 20      5      1655.864  1342.242
## 21      1      1570.350  1672.057
## 22      2      1505.834  1931.884
## 23      3      1597.794  1562.194
## 24      4      1624.176  1456.262
## 25      5      1652.975  1344.246
## 26      1      1569.490  1673.654
## 27      2      1505.115  1932.817
## 28      3      1597.749  1561.237
## 29      4      1623.832  1455.656
## 30      5      1652.459  1344.168
## 31      1      1567.805  1673.954
## 32      2      1504.501  1929.448
## 33      3      1596.183  1560.627
## 34      4      1623.032  1452.553
## 35      5      1650.304  1346.741
## 36      1      1566.402  1675.732
## 37      2      1503.501  1929.308
## 38      3      1595.927  1558.254
## 39      4      1622.292  1451.201
## 40      5      1649.215  1346.675
## 41      1      1566.388  1676.108
## 42      2      1503.474  1929.203
## 43      3      1594.927  1564.856
## 44      4      1621.855  1456.666
## 45      5      1648.872  1348.656
## 46      1      1564.032  1673.956
## 47      2      1501.270  1926.369
## 48      3      1592.954  1560.655
## 49      4      1619.690  1453.864
## 50      5      1645.779  1350.076
## 51      1      1563.755  1676.266
## 52      2      1501.270  1926.371
## 53      3      1592.941  1560.768
## 54      4      1619.586  1454.893
## 55      5      1645.736  1350.627
```

```
degree_0 <- colMeans(ldply(1:5, .fun = do.chunk_2, chunkdef = cv, dat = Wage, a= 0))
degree_1 <- colMeans(ldply(1:5, .fun = do.chunk_2, chunkdef = cv, dat = Wage, a= 1))
degree_2 <- colMeans(ldply(1:5, .fun = do.chunk_2, chunkdef = cv, dat = Wage, a= 2))
degree_3 <- colMeans(ldply(1:5, .fun = do.chunk_2, chunkdef = cv, dat = Wage, a= 3))
degree_4 <- colMeans(ldply(1:5, .fun = do.chunk_2, chunkdef = cv, dat = Wage, a= 4))
degree_5 <- colMeans(ldply(1:5, .fun = do.chunk_2, chunkdef = cv, dat = Wage, a= 5))
degree_6 <- colMeans(ldply(1:5, .fun = do.chunk_2, chunkdef = cv, dat = Wage, a= 6))
```

```

degree_7 <- colMeans(ldply(1:5, .fun = do.chunk_2, chunkdef = cv, dat = Wage, a= 7))
degree_8 <- colMeans(ldply(1:5, .fun = do.chunk_2, chunkdef = cv, dat = Wage, a= 8))
degree_9 <- colMeans(ldply(1:5, .fun = do.chunk_2, chunkdef = cv, dat = Wage, a= 9))
degree_10 <- colMeans(ldply(1:5, .fun = do.chunk_2, chunkdef = cv, dat = Wage, a= 10))

```

```

df2 = as.data.frame(rbind(degree_0, degree_1, degree_2, degree_3, degree_4, degree_5, degree_6, degree_7, degree_8, degree_9, degree_10))
df2$degree = 0:10
df2[-1]

```

##	train.error	val.error	degree
## degree_0	1740.601	1741.540	0
## degree_1	1673.842	1676.139	1
## degree_2	1597.556	1600.107	2
## degree_3	1592.273	1595.125	3
## degree_4	1590.226	1593.329	4
## degree_5	1589.729	1593.506	5
## degree_6	1588.365	1592.664	6
## degree_7	1587.467	1592.234	7
## degree_8	1587.103	1595.098	8
## degree_9	1584.745	1592.984	9
## degree_10	1584.658	1593.785	10


(c):

```

df2 %>% ggplot() +
  geom_line(aes(x = degree, y = train.error, color = "red"), na.rm = TRUE) +
  geom_line(aes(x = degree, y = val.error, color = "blue"))+
  scale_color_discrete("legend", labels = c("train.error", "test.error")) +
  labs(title = "Training and Testing Error of wages as a polynomial function of of age", x = "Degree of

```

Training and Testing Error of wages as a polynomial function of of age



As p increases the training error and the testing error both goes significantly down after 1. However we can see that the training error has a higher error than the test error. At around Degree 10 we can see that it has the minimum test error. So we choose the model at degree 10 because we want to minimize the test errors.